

Optimization

Jungdam Won

Computer Science & Engineering
Seoul National Univ.



Agenda

- What is optimization?
 - Terminologies
- Methods to solve the optimization problem
 - Analytical methods
 - Numerical methods
- Some examples and special types of problems

What is Optimization?



- The selection of a best element, with regard to some criterion, from some set of available alternatives. -Wikipedia-
- Examples
 - How to design a meal with enough nutrition given limited budget?
 - What is the best VLSI (IC chip) design to minimize its size while satisfying the performance criteria?
 - What are the parameters of deep neural networks that provide the best classification performance?

What is Optimization?





Find a **minimum value** of $f(x_1, x_2, \dots, x_n)$
among all possible combination of variables x_1, x_2, \dots, x_n

What is Optimization?



Find a **set of variables** (x_1, x_2, \dots, x_n) that give the minimum value of $f(x_1, x_2, \dots, x_n)$

What is Optimization?



Find a **set of variables** (x_1, x_2, \dots, x_n) that give the minimum value of $f(x_1, x_2, \dots, x_n)$ among the variables satisfying the constraints

What is Optimization?

- A minimization problem can be replaced by a maximization, and vice versa


$$\operatorname{argmin}_{x_1, x_2, \dots, x_n} f(x_1, x_2, \dots, x_n) \leftrightarrow \operatorname{argmax}_{x_1, x_2, \dots, x_n} -f(x_1, x_2, \dots, x_n)$$

- Constants can be ignored while solving the problem



$$\operatorname{argmin}_{x_1, x_2, \dots, x_n} f(x_1, x_2, \dots, x_n) + c \leftrightarrow \operatorname{argmin}_{x_1, x_2, \dots, x_n} f(x_1, x_2, \dots, x_n)$$

Minima

- A real value function f defined on X has a global minima x^* if $f(x^*) \leq f(x)$ for all $x \in X$
- A real value function f defined on X has a local minima if there exists ϵ such that $f(x^*) \leq f(x)$ for all x within distance ϵ of x^*




Minima




- In general, there could exist many minima/maxima given an objective function and constraints except for some special cases like the figure on the right

Simple Examples


- $\min_x [x^2 + 0.1]$



- $\operatorname{argmax}_{x_1, x_2} [-(x_1^2 + x_2^2) + 4]$



- $\operatorname{argmax}_x [e^{-(x-1)^2}]$




How to solve the problems?

- Analytical Methods
 - Unconstrained case
 - Lagrange Multipliers
 - KKT conditions
- Numerical Methods
 - Grid Search (brute-force)
 - Gradient-based Methods
 - Gradient-free Methods

Analytical Methods: An Example

$$\operatorname{argmin}_x f(x) = \operatorname{argmin}_x (x - 1)^2$$



Analytical Methods: An Example

$$\operatorname{argmin}_x f(x) = \operatorname{argmin}_x (x - 1)^2$$

$$\frac{df(x)}{dx} = \frac{d(x - 1)^2}{dx} = 2(x - 1) = 0$$

$$\therefore x = 1$$

- Geometrical meaning of critical points of which gradients are zero
 - Landscape of surrounding area looks flat
- But, this is a necessary condition to become a minimum and maximum

Gradient (First-order Derivative)

- The gradient of a scalar-valued differentiable function f of multiple variables is the vector-valued function $\nabla f(\mathbf{x})$ whose value at a point \mathbf{x} is the "direction and rate of fastest increase". - Wiki -

$$\nabla f(\mathbf{x}) = \nabla_{\mathbf{x}} f = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

Hessian (Second-order Derivative)

- The Hessian is a square matrix of second-order partial derivatives of a scalar-valued function. It describes the local curvature of a function of many variables. - Wiki -

$$\nabla^2 f(\mathbf{x}) = \nabla_{\mathbf{xx}}^2 f = H_f(\mathbf{x}) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

Analytical Methods: An Example

$$\underset{x_1, x_2}{\operatorname{argmin}} [(x_1 - 1)^2 + 3(x_2 - 2)^2]$$


$$= \underset{x_1, x_2}{\operatorname{argmin}} \left[[x_1 \quad x_2] \begin{bmatrix} 1 & 0 \\ 0 & 3 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + [-2 \quad -12] \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + 13 \right]$$

$$\nabla f(x) = \begin{bmatrix} \frac{\partial f(x)}{\partial x_1} \\ \frac{\partial f(x)}{\partial x_2} \end{bmatrix} = \begin{bmatrix} 2(x_1 - 1) \\ 6(x_2 - 2) \end{bmatrix} = \mathbf{0}$$

$$\therefore x_1 = 1, x_2 = 2$$


Isocontours

$$\underset{x_1, x_2}{\operatorname{argmin}} [(x_1 - 1)^2 + 3(x_2 - 2)^2]$$




Equality Constraints

$$\begin{aligned} & \underset{x_1, x_2}{\operatorname{argmin}} [(x_1 - 1)^2 + 3(x_2 - 2)^2] \\ & \text{s.t. } x_1 - 2x_2 = 0 \end{aligned}$$





Equality Constraints

$$\underset{x_1, x_2}{\operatorname{argmax}}[x_1 + x_2] \\ \text{s.t. } x_1^2 + x_2^2 - 4 = 0$$



Equality Constraints



- At the minimum/maximum, the gradient of the objective function and the gradient of equality constraint are parallel

Lagrange Multipliers

- There should exist λ (Lagrange multiplier) which makes $\nabla f(\mathbf{x})$ and $\nabla c(\mathbf{x})$ the same

$$\nabla f(\mathbf{x}) - \lambda \nabla c(\mathbf{x}) = 0$$

- Similar statement holds true for the multiple constraints, refer to the doc below for the proof
 - <https://personal.math.ubc.ca/~feldman/m226/multiLagrange.pdf>

$$\nabla f(\mathbf{x}) - \sum \lambda_i \nabla c_i (\mathbf{x}) = 0$$

Lagrange Multipliers

- To find the solution the optimization problem with equality constraints

$$\underset{\mathbf{x}}{\operatorname{argmax}} f(\mathbf{x}) \\ \text{subject to } c(\mathbf{x}) = 0$$

- Form a Lagrange function $L(\mathbf{x}, \lambda)$ then set its all partial derivatives (also including λ) as zero

$$L(\mathbf{x}, \lambda) = f(\mathbf{x}) - \lambda c(\mathbf{x})$$

$$\nabla_{\mathbf{x}} L(\mathbf{x}, \lambda) = 0 \rightarrow \nabla f(\mathbf{x}) = \lambda \nabla c(\mathbf{x})$$

$$\nabla_{\lambda} L(\mathbf{x}, \lambda) = 0 \rightarrow c(\mathbf{x}) = 0$$

Lagrange Multipliers

$$\underset{x_1, x_2}{\operatorname{argmax}} [x_1 + x_2] \text{ s.t. } x_1^2 + x_2^2 - 4 = 0$$

$$L(x_1, x_2, \lambda) = x_1 + x_2 - \lambda(x_1^2 + x_2^2 - 4)$$

$$\begin{aligned} \nabla_{x_1} &= 1 - 2\lambda x_1 = 0 \\ \nabla_{x_2} &= 1 - 2\lambda x_2 = 0 \end{aligned} \quad \left. \right\} \quad x_1 = x_2$$

$$\nabla_{\lambda} = x_1^2 + x_2^2 - 4 = 0$$

$$\therefore x_1 = x_2 = \sqrt{2}$$

Lagrange Multipliers

$$\begin{aligned} & \underset{x_1, x_2}{\operatorname{argmax}} [(x_1 - 1)^2 + 3(x_2 - 2)^2] \\ & \text{s.t. } x_1 - 2x_2 = 0 \end{aligned}$$

$$\therefore x_1 = \frac{16}{7}, x_2 = \frac{8}{7}$$

Lagrange Multipliers

$$\begin{aligned} & \underset{x_1, x_2}{\operatorname{argmax}} [(x_1 - 1)^2 + 3(x_2 - 2)^2] \\ & \text{s.t. } x_1 - 2x_2 = 0 \end{aligned}$$

$$L(x_1, x_2, \lambda) = (x_1 - 1)^2 + 3(x_2 - 2)^2 + \lambda(x_1 - 2x_2)$$

$$\left. \begin{array}{l} \nabla_{x_1} = 2(x_1 - 1) + \lambda = 0 \\ \nabla_{x_2} = 6(x_2 - 2) - 2\lambda = 0 \end{array} \right\} 2x_1 + 3x_2 = 8$$

$$\nabla_{\lambda} = x_1 - 2x_2 = 0$$

$$\therefore x_1 = \frac{16}{7}, x_2 = \frac{8}{7}$$

Inequality Constraints

- The optimization problem with inequality constraints is harder to solve in general than the optimization with no constraint or equality constraints
- **Karush–Kuhn–Tucker (KKT)** conditions can be used to solve general problems but many problems with simple forms can be solved more efficiently

Linear Programming

- Linear Programming (LP) is a special case of the optimization problem where its objective function and constraints are linear

$$\begin{aligned} & \underset{\mathbf{x}}{\operatorname{argmin}} \mathbf{c}^T \mathbf{x} \\ \text{s. t. } & A \mathbf{x} \leq \mathbf{b} \\ & \mathbf{x} \geq 0 \end{aligned}$$

Linear Programming: An Example


$$\underset{x_1, x_2}{\operatorname{argmax}} \quad 2x_1 + x_2$$

$$\text{s.t. } x_1 + x_2 \leq 3$$

$$x_1 - x_2 \leq 2$$

$$x_1 \geq 0$$

$$x_2 \geq 0$$



- In LP, the solution is located at the *boundaries*


Quadratic Programming

- Quadratic Programming (QP) is a special case of the optimization problem where its objective function is quadratic and constraints are linear inequalities

$$\begin{aligned} & \underset{\mathbf{x}}{\operatorname{argmin}} \frac{1}{2} \mathbf{x}^T \mathbf{Q} \mathbf{x} + \mathbf{c}^T \mathbf{x} \\ & \text{s.t. } \mathbf{A} \mathbf{x} \leq \mathbf{b} \\ & \quad \mathbf{x} \geq 0 \end{aligned}$$

Quadratic Programming: An Example

$$\begin{aligned} & \underset{x}{\operatorname{argmax}} (x - 1)^2 \\ \text{s.t. } & x_1 - 2x_2 \leq 0 \end{aligned}$$



Quadratic Programming: An Example

- QP with linear equality constraints can be solved efficiently


$$\begin{aligned} & \underset{\mathbf{x}}{\operatorname{argmin}} \frac{1}{2} \mathbf{x}^T \mathbf{Q} \mathbf{x} - \mathbf{c}^T \mathbf{x} \\ & \text{ s. t. } \mathbf{A} \mathbf{x} - \mathbf{b} = \mathbf{0} \end{aligned}$$

$$\begin{bmatrix} \mathbf{Q} & \mathbf{A}^T \\ \mathbf{A} & \mathbf{0} \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \boldsymbol{\lambda} \end{bmatrix} = \begin{bmatrix} \mathbf{c} \\ \mathbf{b} \end{bmatrix} \quad : \text{KKT condition}$$

$$\therefore \begin{bmatrix} \mathbf{x} \\ \boldsymbol{\lambda} \end{bmatrix} = \begin{bmatrix} \mathbf{Q} & \mathbf{A}^T \\ \mathbf{A} & \mathbf{0} \end{bmatrix}^{-1} \begin{bmatrix} \mathbf{c} \\ \mathbf{b} \end{bmatrix}$$

Convexity

- To understand the type of a critical point \mathbf{x} ($\nabla f(\mathbf{x}) = 0$), the second derivative (Hessian) can be used
 - Positive: Minima / Negative: Maxima / Zero: Saddle Pt.



Convexity



- Given a critical point x satisfies given constraints, if the function is **convex**, then it is guaranteed to be **the global minimum**
- A real-valued function is called **convex** if the line segment between any two points on the graph of the function lies **above** the graph between the two points.

Convexity

- Given a convex set X and a real valued function $f: X \rightarrow R$, the function f is **convex** if and only if

$$f(tx_1 + (1 - t)x_2) \leq tf(x_1) + (1 - t)f(x_2)$$

for all $0 \leq t \leq 1$ and $x_1, x_2 \in X$



Numerical Methods


- Up to now, we learned how to solve the optimization problem analytically
- However, many practical problems can't be easily formulated as a combination of simple equations which can be solved by hand
 - E.g., 1000 optimization variables
- Instead, we develop an algorithm (procedure) to find one of the local minima

Numerical Methods

- Grid Search (brute-force)
- Gradient-based Methods
- Gradient-free Methods

Grid Search (Brute-Force)

$$\min_x f(x) = \max_x [e^{-(x-1)^2}]$$



Grid Search (Brute-Force)


- Pros
 - Guaranteed to find the minimum (assuming fine enough grid)
 - Simple and easy to implement
- Cons
 - Computationally inefficient
 - Only works for small number of parameters
 - 5 parameters and 100 slices => 10^{10} evaluations are needed

Gradient-based Methods

- First-order vs. Second-order methods
 - The first order method only uses **gradient** (the first derivative) information only
 - The second order method uses **gradient** (the first derivative) and **Hessian** (the second derivative) informationType equation here.


Gradient Descent

$$\min_x f(x) = \min_x (x - 1)^2$$




Gradient Descent

$$\min_x f(x) = \min_x (x - 1)^2$$




Gradient Descent

$$\min_x f(x) = \min_x (x - 1)^2$$



Gradient Descent


$$\min_x f(x) = \min_x (x - 1)^2$$



Gradient Descent

- It is a first-order iterative optimization algorithm for finding a local minimum of a differentiable function. The idea is to take repeated steps in the opposite direction of the gradient (or approximate gradient) of the function at the current point - Wikipedia -

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \gamma \nabla f(\mathbf{x}_t)$$



Gradient Descent

- What if it is hard to calculate the gradient directly?
- We could use approximation
 - But, it requires extra number of evaluations on the objective function


$$\nabla f(\mathbf{x}_t) \approx \frac{\nabla f(\mathbf{x}_t + \epsilon) - \nabla f(\mathbf{x}_t - \epsilon)}{2\epsilon}$$

Problems with Gradient Descent

- If the step-size is too small, gradient descent can be very slow
- If the step-size is too large, it can easily overshoot, fail to converge, or even diverge
- Choosing an appropriate step-size is important but not trivial
 - There exist many algorithms for how to change the step-size adaptively


Sequential Quadratic Programming

- Sequential quadratic programming (SQP) finds a solution by iteratively solving a quadratic programming via linearizing the objective and constraint functions at each timestep




Sequential Quadratic Programming

- Sequential quadratic programming (SQP) finds a solution by iteratively solving a quadratic programming via linearizing the objective and constraint functions at each timestep




Sequential Quadratic Programming

- Sequential quadratic programming (SQP) finds a solution by iteratively solving a quadratic programming via linearizing the objective and constraint functions at each timestep




Sequential Quadratic Programming

- Sequential quadratic programming (SQP) finds a solution by iteratively solving a quadratic programming via linearizing the objective and constraint functions at each timestep



Sequential Quadratic Programming

- Sequential quadratic programming (SQP) finds a solution by iteratively solving a quadratic programming via linearizing the objective and constraint functions at each timestep



Sequential Quadratic Programming

$$\begin{aligned} & \underset{\mathbf{x}}{\operatorname{argmin}} f(\mathbf{x}) \\ \text{subject to } & b(\mathbf{x}_t) \leq \mathbf{0} \\ & c(\mathbf{x}_t) = \mathbf{0} \end{aligned}$$

Initialize \mathbf{x}_t

Until $|\mathbf{x}_{t+1} - \mathbf{x}_t| < \epsilon$

$$\begin{aligned} \mathbf{d} = \underset{\mathbf{d}}{\operatorname{argmin}} & [f(\mathbf{x}_t) + \nabla f(\mathbf{x}_t)^T \mathbf{d} + \frac{1}{2} \mathbf{d}^T \nabla^2 f(\mathbf{x}_t) \mathbf{d}] \\ \text{s. t. } & b(\mathbf{x}_t) + \nabla b(\mathbf{x}_t)^T \mathbf{d} \leq \mathbf{0} \\ & c(\mathbf{x}_t) + \nabla c(\mathbf{x}_t)^T \mathbf{d} = \mathbf{0} \end{aligned}$$

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \gamma \mathbf{d}$$

Constraints


- Projection Method
 - Solve QP with equality constraints only
 - If some variables do not satisfy the inequality constraints, perform projection
- Transfer from hard-constraints to soft-constraints
 - Reformulate the problem so that the constraints are included in the objective function

Gradient-free (Blackbox) Methods

- Particle Swarm
- Simulated Annealing
- Covariance Matrix Adaptation Evolution Strategy
- Genetic Algorithm

Particle Swarm Optimization

- PSO creates a set of solutions (**particles**) with velocities and iteratively moves these particles around in the search-space some criteria



Simulated Annealing


- An optimization algorithm that imitates the annealing process in metallurgy

Algorithm 1: Simulated Annealing (for maximization)

Initialize a solution x ;
for $t = T_{max}$ to T_{min} **do**
 $f(x) \leftarrow$ Evaluate the current solution;
 $x' \leftarrow$ Select a new solution;
 $f(x') \leftarrow$ Evaluate the new solution;
 $\Delta x = f(x'_i) - f(x_i)$;
 $\rho = \frac{1}{1+exp(-\Delta x/t)}$;
 $x \leftarrow$ Accept a new solution with the probability ρ
end

CMA-ES

- Covariance Matrix Adaptation Evolution Strategy (CMA-ES) updates a stochastic distribution based on evolution strategy



Genetic Algorithm

- A method inspired by the process of natural selection process, where biologically inspired operators such as mutation, crossover, and selection are used to refine candidate solutions

Algorithm 1: GeneticAlgorithm

```
Initialize random population;  
for  $g = 1$  to  $g_{max}$  do  
    Select parents from population;  
    Perform reproduction (of children) from the selected parents;  
    Perform crossover with some probability;  
    Perform mutation with some probability;  
    Reduce the population based on the evaluation ;  
end
```

Summary

- We learned some fundamental concepts of the optimization problem, and some popular methods to solve the problem
- Nowadays, many efficient algorithms are readily available, so it is more important to understand which optimization solver is appropriate to your problems than you can implement specific algorithms unless you are a researcher who develop an algorithm

